

```

### v2.1

# Importing the necessary packages

import tensorflow as tf
import numpy as np
import scipy.misc
from tensorflow.keras.applications.resnet_v2 import ResNet50V2
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.resnet_v2 import preprocess_input,
decode_predictions
from tensorflow.keras import layers
from tensorflow.keras.layers import Input, Add, Dense, Activation,
ZeroPadding2D, Flatten, Conv2D, AveragePooling2D, MaxPooling2D,
GlobalMaxPooling2D
from tensorflow.keras.models import Model, load_model
from resnets_utils import *
from tensorflow.keras.initializers import random_uniform,
glorot_uniform, constant, identity
from tensorflow.python.framework.ops import EagerTensor
from matplotlib.pyplot import imshow

from test_utils import summary, comparator
import public_tests

%matplotlib inline
np.random.seed(1)
tf.random.set_seed(2)

# The problem of deep neural networks is vanishing/exploding
gradients.

# Note
# In the case of deep neural networks, there is the problem of
vanishing gradients.
# The gradients of the initial layer weights are very small as
compared to the gradients of the weights of layers at the back.
# So Adam's optimization considers adaptive learning rate to be large
for the initial layers and small for the layers at
# the back. Still the deep neural networks suffer from vanishing
gradients.

# In the rare case of exploding gradients, the optimizing algorithm
may diverge.

# Due to vanishing gradients, the learning speed of weights of initial
layers is small when compared to the learning speed
# of weights of layers at the back.

```

```

# Resnets make it easier for the model to learn the identity function.
# The gradient of the identity function is 1.
# So this reduces the chances of vanishing/exploding gradients which
# happens as we backpropagate (multiplication of matrices)
# from the back to the initial layers.

# Two main types of blocks are used in a ResNet, depending mainly on
# whether the input/output dimensions are the same or
# different. We are going to implement both of them: the "identity
# block" and the "convolutional block."

# Here is where you're actually using the power of the Functional API
# to create a shortcut path

```

## Identity block

```

def identity_block(X, f, filters, initializer=random_uniform):
    """
    Implementation of the identity block as defined in Figure 4

    Arguments:
    X -- input tensor of shape (m, n_H_prev, n_W_prev, n_C_prev)
    f -- integer, specifying the shape of the middle CONV's window for
    the main path
    filters -- python list of integers, defining the number of filters
    in the CONV layers of the main path
    initializer -- to set up the initial weights of a layer. Equals to
    random uniform initializer

    Returns:
    X -- output of the identity block, tensor of shape (m, n_H, n_W,
    n_C)
    """

    # Retrieve Filters
    F1, F2, F3 = filters

    # Save the input value
    X_shortcut = X
    # 'X_shortcut' and 'X' are of shape nh x nw x nc
    # Note - Tensors are immutable unlike python lists and arrays

    # First component of main path
    X = Conv2D(filters = F1, kernel_size = 1, strides = (1,1), padding
    = 'valid', kernel_initializer = initializer(seed=0))(X)
    # 'X' is of shape nh x nw x F1
    # kernel_initializer refers to the method used for the initial
    weight initialization of the kernel (or weights) in a
    # layer of a neural network.
    # 'valid' padding means no padding

```

```

X = BatchNormalization(axis = 3)(X) # Default axis
X = Activation('relu')(X)

## Second component of main path
## Set the padding = 'same'
X = Conv2D(filters = F2, kernel_size = f, strides = (1,1), padding
= 'same', kernel_initializer = initializer(seed=0))(X)
# 'X' is of shape nh x nw x F2
X = BatchNormalization(axis = 3)(X)
X = Activation('relu')(X)

## Third component of main path
## Set the padding = 'valid'
X = Conv2D(filters = F3, kernel_size = 1, strides = (1,1), padding
= 'valid', kernel_initializer = initializer(seed=0))(X)
# 'X' is of shape nh x nw x F3
X = BatchNormalization(axis = 3)(X)

## Final step: Add shortcut value to main path, and pass it
through a RELU activation
X = Add()([X,X_shortcut])
# Since this is an identity block, F3 = nc because X and
X_shortcut have to be the same size
X = Activation('relu')(X)

return X

### you cannot edit this cell

tf.keras.backend.set_learning_phase(False)

np.random.seed(1)
tf.random.set_seed(2)
X1 = np.ones((1, 4, 4, 3)) * -1
X2 = np.ones((1, 4, 4, 3)) * 1
X3 = np.ones((1, 4, 4, 3)) * 3

X = np.concatenate((X1, X2, X3), axis = 0).astype(np.float32)

A3 = identity_block(X, f=2, filters=[4, 4, 3],
                    initializer=lambda seed=0:constant(value=1))
print('\033[1mWith training=False\033[0m\n')
A3np = A3.numpy()
print(np.around(A3.numpy()[:(0,-1),:,:].mean(axis = 3), 5))
resume = A3np[:,(0,-1),:,:].mean(axis = 3)
print(resume[1, 1, 0])

tf.keras.backend.set_learning_phase(True)

print('\n\033[1mWith training=True\033[0m\n')

```

```

np.random.seed(1)
tf.random.set_seed(2)
A4 = identity_block(X, f=2, filters=[3, 3, 3],
                    initializer=lambda seed=0:constant(value=1))
print(np.around(A4.numpy()[:(0,-1),:,:].mean(axis = 3), 5))

public_tests.identity_block_test(identity_block)

With training=False

[[[ 0.      0.      0.      0.      ]
  [ 0.      0.      0.      0.      ]]]

[[192.99992 192.99992 192.99992 96.99996]
 [ 96.99996 96.99996 96.99996 48.99998]]

[[578.99976 578.99976 578.99976 290.99988]
 [290.99988 290.99988 290.99988 146.99994]]]
96.99996

With training=True

[[[0.      0.      0.      0.      ]
  [0.      0.      0.      0.      ]]]

[[0.40732 0.40732 0.40732 0.40732]
 [0.40732 0.40732 0.40732 0.40732]]

[[5.00011 5.00011 5.00011 3.25955]
 [3.25955 3.25955 3.25955 2.40732]]]
All tests passed!

# Note - The CONV2D layer on the shortcut path does not use any non-
linear activation function.
# Its main role is to just apply a (learned) linear function that
reduces the dimension of the input,
# so that the dimensions match up for the later addition step.

```

## Convolutional block

```

def convolutional_block(X, f, filters, s = 2,
initializer=glorot_uniform):
    """
    Implementation of the convolutional block as defined in Figure 4

    Arguments:
    X -- input tensor of shape (m, n_H_prev, n_W_prev, n_C_prev)
    f -- integer, specifying the shape of the middle CONV's window for
the main path
    filters -- python list of integers, defining the number of filters

```

*in the CONV layers of the main path*  
*s -- Integer, specifying the stride to be used*  
*initializer -- to set up the initial weights of a layer. Equals to*  
*Glorot uniform initializer,*  
*also called Xavier uniform initializer.*

*Returns:*

*X -- output of the convolutional block, tensor of shape (m, n\_H,*  
*n\_W, n\_C)*  
"""

*# Retrieve Filters*

F1, F2, F3 = filters

*# Save the input value*

X\_shortcut = X

##### MAIN PATH #####

*# First component of main path*

X = Conv2D(filters = F1, kernel\_size = 1, strides = (s, s),  
padding='valid', kernel\_initializer = initializer(seed=0))(X)

X = BatchNormalization(axis = 3)(X)

X = Activation('relu')(X)

*# X is of shape ((nh-1)/2) + 1, ((nw-1)/2) + 1, F1*

*## Second component of main path*

X = Conv2D(filters = F2, kernel\_size = f, strides = 1,  
padding='same', kernel\_initializer = initializer(seed=0))(X)

X = BatchNormalization(axis = 3)(X)

X = Activation('relu')(X)

*# X is of shape ((nh-1)/2) + 1, ((nw-1)/2) + 1, F2*

*## Third component of main path*

X = Conv2D(filters = F3, kernel\_size = 1, strides = 1,  
padding='valid', kernel\_initializer = initializer(seed=0))(X)

X = BatchNormalization(axis = 3)(X)

*# X is of shape ((nh-1)/2) + 1, ((nw-1)/2) + 1, F3*

##### SHORTCUT PATH #####

X\_shortcut = Conv2D(filters = F3, kernel\_size = 1, strides = (s,  
s), padding='valid', kernel\_initializer = initializer(seed=0))  
(X\_shortcut)

X\_shortcut = BatchNormalization(axis = 3)(X\_shortcut)

*# X\_shortcut is of shape ((nh-1)/2) + 1, ((nw-1)/2) + 1, F3*

*# Add shortcut value to main path and pass it through a RELU*  
*activation*

X = Add()([X, X\_shortcut])

```

X = Activation('relu')(X)

return X

# Note

# In ResNet, there are 2 blocks - identity blocks and convolutional
# blocks. I can understand how identity blocks result in the
# identity function. But in the case of convolutional blocks,
# X_shortcut is also passed through 1 convolutional layer.
# But there is no activation function and hence there is no non-linear
# transformation and this kind of represents the
# identity function. Because there is only linear transformation.

# This is not strictly identity function but it facilitates the
# training of deeper networks by maintaining stronger
# gradient flows, even when the dimensions change.

### you cannot edit this cell

public_tests.convolutional_block_test(convolutional_block)

tf.Tensor(
[[[0.33485505 1.6415989 0.33789736 0.08511472 0.814965 0.
    [0.17509979 1.5699672 0.2606045 0. 0.767209 0.
    ]

    [[0. 1.4983511 0.16896994 0. 0.61830646 0.
    [0. 1.4502985 0.11632714 0. 0.58068544
    0.
    ]]], shape=(2, 2, 6), dtype=float32)
All tests passed!

```

## Building of ResNet-50 model

```

# Details of ResNet-50 model

# Zero-padding pads the input with a pad of (3,3)
# Stage 1:
# The 2D Convolution has 64 filters of shape (7,7) and uses a stride
# of (2,2).
# BatchNorm is applied to the 'channels' axis of the input.
# MaxPooling uses a (3,3) window and a (2,2) stride.

# Stage 2:
# The convolutional block uses three sets of filters of size
# [64,64,256], "f" is 3, and "s" is 1.
# The 2 identity blocks use three sets of filters of size [64,64,256],
# and "f" is 3.

# Stage 3:
# The convolutional block uses three sets of filters of size
# [128,128,512], "f" is 3 and "s" is 2.

```

```

# The 3 identity blocks use three sets of filters of size
[128,128,512] and "f" is 3.

# Stage 4:
# The convolutional block uses three sets of filters of size [256,
256, 1024], "f" is 3 and "s" is 2.
# The 5 identity blocks use three sets of filters of size [256, 256,
1024] and "f" is 3.

# Stage 5:
# The convolutional block uses three sets of filters of size [512,
512, 2048], "f" is 3 and "s" is 2.
# The 2 identity blocks use three sets of filters of size [512, 512,
2048] and "f" is 3.
# The 2D Average Pooling uses a window of shape (2,2).
# The 'flatten' layer doesn't have any hyperparameters.
# The Fully Connected (Dense) layer reduces its input to the number of
classes using a softmax activation.

# Calculation of total number of layers in ResNet-50 model

# 12 Identity_blocks = 36 layers
# 4 convolutional_blocks = 12 layers
# 1 convolutional_layer = 1 layer
# 1 fully connected layer = 1 layer

# Total = 50 layers

def ResNet50(input_shape = (64, 64, 3), classes = 6, training=False):
    """
    Stage-wise implementation of the architecture of the popular
    ResNet50:
    CONV2D -> BATCHNORM -> RELU -> MAXPOOL -> CONVBLOCK -> IDBLOCK*2 -
    > CONVBLOCK -> IDBLOCK*3
    -> CONVBLOCK -> IDBLOCK*5 -> CONVBLOCK -> IDBLOCK*2 -> AVGPPOOL ->
    FLATTEN -> DENSE

    Arguments:
    input_shape -- shape of the images of the dataset
    classes -- integer, number of classes

    Returns:
    model -- a Model() instance in Keras
    """

    # Define the input as a tensor with shape input_shape
    X_input = Input(input_shape)

    # Zero-Padding

```

```

X = ZeroPadding2D((3, 3))(X_input)

# Stage 1
X = Conv2D(64, (7, 7), strides = (2, 2), kernel_initializer =
glorot_uniform(seed=0))(X)
X = BatchNormalization(axis = 3)(X)
X = Activation('relu')(X)
X = MaxPooling2D((3, 3), strides=(2, 2))(X)

# Stage 2
X = convolutional_block(X, f = 3, filters = [64, 64, 256], s = 1)
# After above step, suppose X is of shape nh,nw,256
X = identity_block(X, 3, [64, 64, 256])
# After above step, X is of shape nh,nw,256
X = identity_block(X, 3, [64, 64, 256])
# After above step, X is of shape nh,nw,256

# Note - For the convolutional block, we can blindly take any
values of f,filters,s
# But for identity_block, if input is of shape (nh,nw,nc), then
filters is [F1,F2,nc]. f can be any value

# Use the instructions above in order to implement all of the
Stages below
# Make sure you don't miss adding any required parameter

## Stage 3
X = convolutional_block(X, f = 3, filters = [128, 128, 512], s =
2)
# After above step, X is of shape ((nh-1)/2) + 1, ((nw-1)/2) + 1,
512 or (nh',nw',512)

# the 3 `identity_block` with correct values of `f` and `filters`
for this stage
X = identity_block(X, 3, [128, 128, 512])
# After above step, X is of shape ((nh-1)/2) + 1, ((nw-1)/2) + 1,
512 or (nh',nw',512)
X = identity_block(X, 3, [128, 128, 512])
# After above step, X is of shape ((nh-1)/2) + 1, ((nw-1)/2) + 1,
512 or (nh',nw',512)
X = identity_block(X, 3, [128, 128, 512])
# After above step, X is of shape ((nh-1)/2) + 1, ((nw-1)/2) + 1,
512 or (nh',nw',512)

# Stage 4
# add `convolutional_block` with correct values of `f`, `filters`
and `s` for this stage
X = convolutional_block(X, f = 3, filters = [256, 256, 1024], s =
2)

```



```

    # After above step, X is of shape ((nh'-1)/2) + 1, ((nw'-1)/2) +
    1, 1024 or (nh'',nw'',1024)

    # the 5 `identity_block` with correct values of `f` and `filters`
    for this stage
    X = identity_block(X, 3, [256, 256, 1024])
    # After above step, X is of shape ((nh'-1)/2) + 1, ((nw'-1)/2) +
    1, 1024 or (nh'',nw'',1024)
    X = identity_block(X, 3, [256, 256, 1024])
    # After above step, X is of shape ((nh'-1)/2) + 1, ((nw'-1)/2) +
    1, 1024 or (nh'',nw'',1024)
    X = identity_block(X, 3, [256, 256, 1024])
    # After above step, X is of shape ((nh'-1)/2) + 1, ((nw'-1)/2) +
    1, 1024 or (nh'',nw'',1024)
    X = identity_block(X, 3, [256, 256, 1024])
    # After above step, X is of shape ((nh'-1)/2) + 1, ((nw'-1)/2) +
    1, 1024 or (nh'',nw'',1024)
    X = identity_block(X, 3, [256, 256, 1024])
    # After above step, X is of shape ((nh'-1)/2) + 1, ((nw'-1)/2) +
    1, 1024 or (nh'',nw'',1024)

    # Stage 5
    # add `convolutional_block` with correct values of `f`, `filters`
    and `s` for this stage
    X = convolutional_block(X, f = 3, filters = [512, 512, 2048], s =
    2)
    # After above step, X is of shape ((nh''-1)/2) + 1, ((nw''-1)/2) +
    1, 2048 or (nh''',nw''',2048)

    # the 2 `identity_block` with correct values of `f` and `filters`
    for this stage
    X = identity_block(X, 3, [512, 512, 2048])
    # After above step, X is of shape ((nh''-1)/2) + 1, ((nw''-1)/2) +
    1, 2048 or (nh''',nw''',2048)
    X = identity_block(X, 3, [512, 512, 2048])
    # After above step, X is of shape ((nh''-1)/2) + 1, ((nw''-1)/2) +
    1, 2048 or (nh''',nw''',2048)

    # AVGP00L
    X = AveragePooling2D((2,2))(X)

    # output layer
    X = Flatten()(X)
    X = Dense(classes, activation='softmax', kernel_initializer =
    glorot_uniform(seed=0))(X)

    # Create model
    model = Model(inputs = X_input, outputs = X)

```

```

    return model

tf.keras.backend.set_learning_phase(True)

# The above line of code explicitly sets it to the learning/training phase which is different from inference/prediction phase.

# Differences between training and prediction phase
# During training, dropout randomly sets input units to 0 at a rate of rate at each step, which helps prevent overfitting, and
# batch normalization normalizes the input using the batch statistics.
# During inference, dropout is not applied, and batch normalization uses the moving average and variance calculated during
# training.

# In TensorFlow 2.x and later,
# When calling model.fit(), TensorFlow automatically sets the learning phase to training mode.
# When using model.evaluate() or model.predict(), TensorFlow sets the learning phase to inference mode.

model = ResNet50(input_shape = (64, 64, 3), classes = 6)
print(model.summary())

### you cannot edit this cell

from outputs import ResNet50_summary

model = ResNet50(input_shape = (64, 64, 3), classes = 6)

comparator(summary(model), ResNet50_summary)

All tests passed!

```

## Compiling the model

```

np.random.seed(1)
tf.random.set_seed(2)
opt = tf.keras.optimizers.Adam(learning_rate=0.00015)
model.compile(optimizer=opt, loss='categorical_crossentropy',
metrics=['accuracy'])

```

## Loading the dataset

```

X_train_orig, Y_train_orig, X_test_orig, Y_test_orig, classes =
load_dataset()

# Normalize image vectors
X_train = X_train_orig / 255.
X_test = X_test_orig / 255.

```

```

# Convert training and test labels to one hot matrices
Y_train = convert_to_one_hot(Y_train_orig, 6).T
Y_test = convert_to_one_hot(Y_test_orig, 6).T

print ("number of training examples = " + str(X_train.shape[0]))
print ("number of test examples = " + str(X_test.shape[0]))
print ("X_train shape: " + str(X_train.shape))
print ("Y_train shape: " + str(Y_train.shape))
print ("X_test shape: " + str(X_test.shape))
print ("Y_test shape: " + str(Y_test.shape))

number of training examples = 1080
number of test examples = 120
X_train shape: (1080, 64, 64, 3)
Y_train shape: (1080, 6)
X_test shape: (120, 64, 64, 3)
Y_test shape: (120, 6)

```

## Training the model

```

model.fit(X_train, Y_train, epochs = 10, batch_size = 32)

Epoch 1/10
34/34 [=====] - 16s 63ms/step - loss: 1.8302
- accuracy: 0.3019
Epoch 2/10
34/34 [=====] - 2s 51ms/step - loss: 1.2657 -
accuracy: 0.5259
Epoch 3/10
34/34 [=====] - 2s 51ms/step - loss: 0.9542 -
accuracy: 0.6583
Epoch 4/10
34/34 [=====] - 2s 51ms/step - loss: 0.6513 -
accuracy: 0.7648
Epoch 5/10
34/34 [=====] - 2s 52ms/step - loss: 0.3943 -
accuracy: 0.8574
Epoch 6/10
34/34 [=====] - 2s 51ms/step - loss: 0.2600 -
accuracy: 0.9009
Epoch 7/10
34/34 [=====] - 2s 51ms/step - loss: 0.3027 -
accuracy: 0.8963
Epoch 8/10
34/34 [=====] - 2s 51ms/step - loss: 0.2490 -
accuracy: 0.9065
Epoch 9/10
34/34 [=====] - 2s 52ms/step - loss: 0.1940 -
accuracy: 0.9343

```

```

Epoch 10/10
34/34 [=====] - 2s 50ms/step - loss: 0.1304 -
accuracy: 0.9528

<keras.callbacks.History at 0x7fe61c8be4f0>

preds = model.evaluate(X_test, Y_test)
print ("Loss = " + str(preds[0]))
print ("Test Accuracy = " + str(preds[1]))

4/4 [=====] - 1s 11ms/step - loss: 0.6611 -
accuracy: 0.7917
Loss = 0.661126971244812
Test Accuracy = 0.7916666865348816

```

## Using a pre\_trained model

```

pre_trained_model = load_model('resnet50.h5')

preds = pre_trained_model.evaluate(X_test, Y_test)
print ("Loss = " + str(preds[0]))
print ("Test Accuracy = " + str(preds[1]))

4/4 [=====] - 1s 9ms/step - loss: 0.1596 -
accuracy: 0.9500
Loss = 0.1595880389213562
Test Accuracy = 0.949999988079071

```

### # Note

```

# Very deep "plain" networks don't work in practice because vanishing
gradients make them hard to train.
# Skip connections help address the Vanishing Gradient problem. They
also make it easy for a ResNet block to learn an
# identity function.
# There are two main types of blocks: The identity block and the
convolutional block.
# Very deep Residual Networks are built by stacking these blocks
together.

```